

In [20]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score
from sklearn.metrics import plot_roc_curve, plot_precision_recall_curve
from sklearn.metrics import classification_report, confusion_matrix
from sklearn import metrics
```

DATA PREPROCESSING

In [21]:

```
header=["loc", "v(g)", "ev(g)", "iv(g)", "n", "v", "l", "d", "i", "e", "b", "t", "IOCode", "IOComment", "
data=pd.read_csv("D://Downloads/Software/Software Dataset/promise5_useful.txt", names=header)
data.head()
```

Out[21]:

	loc	v(g)	ev(g)	iv(g)	n	v	l	d	i	e	...	IOCode	IOComi
0	91	9	3	2	318	2089.21	0.04	27.68	75.47	57833.24	...	80	
1	109	21	5	18	381	2547.56	0.04	28.37	89.79	72282.68	...	97	
2	505	106	41	82	2339	20696.93	0.01	75.93	272.58	1571506.88	...	457	
3	107	25	7	14	619	4282.78	0.02	52.91	80.95	226588.75	...	103	
4	74	11	1	8	294	1917.93	0.03	28.77	66.66	55178.46	...	60	

5 rows × 22 columns

In [22]:

```
data=pd.DataFrame(data)

data.defects=data.defects.replace(True,1)
data.defects=data.defects.replace(False,0)
```

In [23]:

```
temp=np.array(data['defects'])
z=0
o=0
for i in temp:
    if(i==1):
        o+=1
    else:
        z+=1
print("ones: %d, zeroes: %d" %(o,z))
```

ones: 76, zeroes: 76

In [24]:

```
arr=np.array(data.defects)
print(np.where(arr==1)) #use shuffle as 1s and 0s are together
```

```
(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15, 16,
        17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33,
        34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50,
        51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67,
        68, 69, 70, 71, 72, 73, 74, 75], dtype=int64),)
```

In [25]:

```
for i in range(16,len(header)-1):
    data[header[i]]=pd.to_numeric(data[header[i]], errors='coerce').astype('float32')
```

In [26]:

```
data=data.dropna(axis=0,how='any')
```

In [27]:

```
defects=data.loc[:, 'defects']
data=data.drop('defects',axis=1)
```

In [28]:

```
from sklearn.preprocessing import Normalizer
transformer=Normalizer().fit(data)
x_scaled=transformer.transform(data)
data = pd.DataFrame(x_scaled,columns = ["loc", "v(g)", "ev(g)", "iv(g)", "n", "v", "l", "d", "i", "e"])
data.head()
```

Out[28]:

	loc	v(g)	ev(g)	iv(g)	n	v	l	d	i
0	0.001570	0.000155	0.000052	0.000035	0.005486	0.036045	6.901112e-07	0.000478	0.001302
1	0.001505	0.000290	0.000069	0.000248	0.005259	0.035167	5.521753e-07	0.000392	0.001239
2	0.000321	0.000067	0.000026	0.000052	0.001486	0.013149	6.352961e-09	0.000048	0.000173
3	0.000471	0.000110	0.000031	0.000062	0.002727	0.018869	8.811351e-08	0.000233	0.000357
4	0.001338	0.000199	0.000018	0.000145	0.005317	0.034683	5.425127e-07	0.000520	0.001205

5 rows × 21 columns

In [29]:

```
data['defects']=defects
#data=data.drop('LOCodeAndComment',axis=1)
#data=data.drop('LOBlank',axis=1)
#data=data.drop('LOComment',axis=1)
data=data.dropna(axis=0,how='any')
data.head()
```


Out[29]:

	loc	v(g)	ev(g)	iv(g)	n	v	l	d	i
0	0.001570	0.000155	0.000052	0.000035	0.005486	0.036045	6.901112e-07	0.000478	0.001302
1	0.001505	0.000290	0.000069	0.000248	0.005259	0.035167	5.521753e-07	0.000392	0.001239
2	0.000321	0.000067	0.000026	0.000052	0.001486	0.013149	6.352961e-09	0.000048	0.000173
3	0.000471	0.000110	0.000031	0.000062	0.002727	0.018869	8.811351e-08	0.000233	0.000357
4	0.001338	0.000199	0.000018	0.000145	0.005317	0.034683	5.425127e-07	0.000520	0.001205

5 rows × 22 columns

In [30]:

```
x=data['loc']
y=data['iv(g)']
z=data['defects']
plt.scatter(x,y,c=z)
plt.show()
```



In [31]:

```
x=data.drop('defects',axis=1).values
y=data[["defects"]].values
```

In [32]:

```
from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest=train_test_split(x,y)
```

In [33]:

```
print(xtrain.shape, ytrain.shape, xtest.shape, ytest.shape)
```

```
(114, 21) (114, 1) (38, 21) (38, 1)
```

In [34]:

```
ytrain, ytest=ytrain.flatten(), ytest.flatten()
```

In [35]:

```
z=0
o=0
for i in ytrain:
    if(i==1):
        o+=1
    else:
        z+=1
print("ones: %d, zeroes: %d" %(o,z))
```

```
ones: 59, zeroes: 55
```

In [36]:

```
for i in range(0,len(ytrain)):
    if(ytrain[i]==1):
        print(xtrain[i])
        print("\n")
```

```
[4.72280452e-03 5.90350566e-04 5.90350566e-04 5.90350566e-04
 2.89271777e-02 1.40527049e-01 8.26490792e-05 4.15016448e-03
 1.99951737e-02 9.87597461e-01 4.72280452e-05 5.48671816e-02
 4.72280452e-03 2.36140226e-03 0.00000000e+00 3.54210339e-03
 6.49385622e-03 1.06263102e-02 1.53491147e-02 1.35780630e-02
 5.90350566e-04]
```

```
[4.71407299e-04 1.10141892e-04 3.08397298e-05 6.16794596e-05
 2.72711325e-03 1.88685397e-02 8.81135138e-08 2.33104301e-04
 3.56639447e-04 9.98276547e-01 6.30011623e-06 5.54597910e-02
 4.53784596e-04 1.40981622e-04 1.76227028e-05 1.71821352e-04
 1.54198649e-04 3.78888109e-04 1.58163757e-03 1.14547568e-03
 1.76227028e-04]
```

```
[2.21651913e-03 5.91071767e-04 4.43303825e-04 2.95535884e-04
 1.27080430e-02 6.76319093e-02 1.03437559e-05 2.17662178e-03
 4.59262763e-03 9.96023901e-01 2.21651913e-05 5.53346612e-02
 2.06875110e-03 7.38830700e-04 1.47767043e-04 1.03437559e-05]
```

BASE PREDICTIONS

1- SVM

In [37]:

```
from sklearn.svm import SVC
```

In [38]:

```
svm_model=SVC()  
svm_model.fit(xtrain,ytrain)
```

Out[38]:

```
SVC(C=1.0, break_ties=False, cache_size=200, class_weight=None, coef0=0.0,  
    decision_function_shape='ovr', degree=3, gamma='scale', kernel='rbf',  
    max_iter=-1, probability=False, random_state=None, shrinking=True,  
    tol=0.001, verbose=False)
```

In [39]:

```
predsvm=svm_model.predict(xtest)  
svm_model.score(xtest,ytest)*100
```

Out[39]:

```
57.89473684210527
```

In [40]:

```

accuracy=confusion_matrix(ytest,predsvm)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predsvm, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predsvm))
print("\n")
print("    P    N")
print(confusion_matrix(ytest,predsvm))
print("\n")
print(classification_report(ytest,predsvm))

```

```

Accuracy:  57.89473684210527
Probability of detection of defect(Recall, pd):  1.0
Probability of false alarm(pf):  0.0
Probability of correct detection(Precision):  0.5151515151515151

```

```

F1-score or FM:  0.6799999999999999
AUC value:  0.6190476190476191

```

```


    P    N
[[ 5 16]
 [ 0 17]]

```

	precision	recall	f1-score	support
0.0	1.00	0.24	0.38	21
1.0	0.52	1.00	0.68	17
accuracy			0.58	38
macro avg	0.76	0.62	0.53	38
weighted avg	0.78	0.58	0.52	38


In [41]:

```
x=np.arange(100)*0.01  
y=x  
  
disp = plot_roc_curve(svm_model, xtest, ytest)  
plt.plot(x,y, '--')  
plt.show()
```



In [42]:

```
disp = plot_precision_recall_curve(svm_model, xtest, ytest)
plt.show()
```



2- KNN

In [116]:

```
from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=11)
```

In [117]:

```
knn.fit(xtrain,ytrain)
```

Out[117]:

```
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=None, n_neighbors=11, p=2,
                    weights='uniform')
```

In [118]:

```
predknn=knn.predict(xtest)
knn.score(xtest,ytest)*100
```

Out[118]:

```
71.05263157894737
```


In [119]:

```

accuracy=confusion_matrix(ytest,predknn)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predknn, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predknn))
print("\n")
print(confusion_matrix(ytest,predknn))
print("\n")
print(classification_report(ytest,predknn))

```

```

Accuracy: 71.05263157894737
Probability of detection of defect(Recall, pd): 0.8235294117647058
Probability of false alarm(pf): 0.1875
Probability of correct detection(Precision): 0.6363636363636364

```

```

F1-score or FM: 0.717948717948718
AUC value: 0.7212885154061625

```

```

[[13  8]
 [ 3 14]]


```

	precision	recall	f1-score	support
0.0	0.81	0.62	0.70	21
1.0	0.64	0.82	0.72	17
accuracy			0.71	38
macro avg	0.72	0.72	0.71	38
weighted avg	0.73	0.71	0.71	38

In [120]:


```
x=np.arange(100)*0.01
y=x

disp = plot_roc_curve(knn, xtest, ytest)
plt.plot(x,y, '--')
plt.show()
```



In [121]:

```
disp = plot_precision_recall_curve(knn, xtest, ytest)
plt.show()
```



In [122]:

```
# try K=1 through K=25 and record testing accuracy
k_range = range(1, 15)

# We can create Python dictionary using [] or dict()
scores = {}

# We use a loop through the range 1 to 26
# We append the scores in the dictionary
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(xtrain, ytrain)
    y_pred = knn.predict(xtest)
    scores[k] = metrics.accuracy_score(ytest, y_pred)

print(scores)
```


```
{1: 0.7105263157894737, 2: 0.7105263157894737, 3: 0.631578947368421, 4: 0.6842105263157895, 5: 0.6842105263157895, 6: 0.6578947368421053, 7: 0.6578947368421053, 8: 0.6578947368421053, 9: 0.631578947368421, 10: 0.6842105263157895, 11: 0.7105263157894737, 12: 0.6842105263157895, 13: 0.6842105263157895, 14: 0.6842105263157895}
```

In [123]:

```
# plot the relationship between K and testing accuracy
# plt.plot(x_axis, y_axis)
plt.plot(k_range, scores)
plt.xlabel('Value of K for KNN')
plt.ylabel('Testing Accuracy')
```

Out[123]:

Text(0, 0.5, 'Testing Accuracy')



3- NAIVE BAYES

In [51]:

```
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
```

In [52]:

```
gnb.fit(xtrain,ytrain)
```

Out[52]:

```
GaussianNB(priors=None, var_smoothing=1e-09)
```

In [53]:

```
predg=gnb.predict(xtest)  
gnb.score(xtest,ytest)*100
```

Out[53]:

```
73.68421052631578
```

In [54]:

```

accuracy=confusion_matrix(ytest,predg)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predg, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predg))
print("\n")
print(confusion_matrix(ytest,predg))
print("\n")
print(classification_report(ytest,predg))

```

```

Accuracy: 73.68421052631578
Probability of detection of defect(Recall, pd): 1.0
Probability of false alarm(pf): 0.0
Probability of correct detection(Precision): 0.6296296296296297

```

```

F1-score or FM: 0.7727272727272727
AUC value: 0.7619047619047619

```

```

[[11 10]
 [ 0 17]]


```

	precision	recall	f1-score	support
0.0	1.00	0.52	0.69	21
1.0	0.63	1.00	0.77	17
accuracy			0.74	38
macro avg	0.81	0.76	0.73	38
weighted avg	0.83	0.74	0.73	38

In [55]:


```
x=np.arange(100)*0.01
y=x

disp = plot_roc_curve(gnb, xtest, ytest)
plt.plot(x,y, '--')
plt.show()
```



In [56]:

```
disp = plot_precision_recall_curve(gnb, xtest, ytest)
plt.show()
```



In [57]:

```
c=0
l=len(ytest)
for i in range(0,l):
    if(predg[i]!=ytest[i]):
        c=c+1
print("Number of mislabeled points out of a total %d points : %d" %(l,c))
```

Number of mislabeled points out of a total 38 points : 10

4- LOGISTIC REGRESSION

In [58]:

```
from sklearn.linear_model import LogisticRegression
logmodel=LogisticRegression()
```

In [59]:

```
logmodel.fit(xtrain,ytrain)
```

Out[59]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                    warm_start=False)
```

In [60]:

```
predlog=logmodel.predict(xtest)
logistic_score=logmodel.score(xtest,ytest)*100
logistic_score
```

Out[60]:

```
60.526315789473685
```

In [61]:

```

accuracy=confusion_matrix(ytest,predlog)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predlog, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predlog))
print("\n")
print(confusion_matrix(ytest,predlog))
print("\n")
print(classification_report(ytest,predlog))

```

```

Accuracy: 60.526315789473685
Probability of detection of defect(Recall, pd): 0.9411764705882353
Probability of false alarm(pf): 0.125
Probability of correct detection(Precision): 0.5333333333333333

```

```

F1-score or FM: 0.6808510638297872
AUC value: 0.6372549019607843

```

```

[[ 7 14]
 [ 1 16]]


```

	precision	recall	f1-score	support
0.0	0.88	0.33	0.48	21
1.0	0.53	0.94	0.68	17
accuracy			0.61	38
macro avg	0.70	0.64	0.58	38
weighted avg	0.72	0.61	0.57	38

In [62]:


```
x=np.arange(100)*0.01
y=x

disp = plot_roc_curve(logmodel, xtest, ytest)
plt.plot(x,y, '--')
plt.show()
```



In [63]:

```
disp = plot_precision_recall_curve(logmodel, xtest, ytest)
plt.show()
```



5- MLP

In [64]:

```
from sklearn.neural_network import MLPClassifier
```

In [65]:

```
model=MLPClassifier(hidden_layer_sizes=(20,20),max_iter=2000)
model.fit(xtrain,ytrain)
```

Out[65]:

```
MLPClassifier(activation='relu', alpha=0.0001, batch_size='auto', beta_1=0.9,
              beta_2=0.999, early_stopping=False, epsilon=1e-08,
              hidden_layer_sizes=(20, 20), learning_rate='constant',
              learning_rate_init=0.001, max_fun=15000, max_iter=2000,
              momentum=0.9, n_iter_no_change=10, nesterovs_momentum=True,
              power_t=0.5, random_state=None, shuffle=True, solver='adam',
              tol=0.0001, validation_fraction=0.1, verbose=False,
              warm_start=False)
```

In [66]:

```
predn=model.predict(xtest)
model.score(xtest,ytest)*100
```

Out[66]:

73.68421052631578

In [67]:

```

accuracy=confusion_matrix(ytest,predn)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predn, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predn))
print("\n")
print(confusion_matrix(ytest,predn))
print("\n")
print(classification_report(ytest,predn))

```

Accuracy: 73.68421052631578
 Probability of detection of defect(Recall, pd): 0.8823529411764706
 Probability of false alarm(pf): 0.13333333333333333
 Probability of correct detection(Precision): 0.6521739130434783

F1-score or FM: 0.75
 AUC value: 0.7507002801120448

```

[[13  8]
 [ 2 15]]


```

	precision	recall	f1-score	support
0.0	0.87	0.62	0.72	21
1.0	0.65	0.88	0.75	17
accuracy			0.74	38
macro avg	0.76	0.75	0.74	38
weighted avg	0.77	0.74	0.73	38

In [68]:


```
x=np.arange(100)*0.01
y=x

disp = plot_roc_curve(model, xtest, ytest)
plt.plot(x,y, '--')
plt.show()
```



In [69]:

```
disp = plot_precision_recall_curve(model, xtest, ytest)
plt.show()
```



6- DECISION TREE

In [70]:

```
from sklearn import tree
```

In [71]:

```
tmodel=tree.DecisionTreeClassifier()  
tmodel.fit(xtrain,ytrain)
```

Out[71]:

```
DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, criterion='gini',  
                      max_depth=None, max_features=None, max_leaf_nodes=None,  
e,  
                      min_impurity_decrease=0.0, min_impurity_split=None,  
                      min_samples_leaf=1, min_samples_split=2,  
                      min_weight_fraction_leaf=0.0, presort='deprecated',  
                      random_state=None, splitter='best')
```

In [72]:

```
predt=tmodel.predict(xtest)  
tmodel.score(xtest,ytest)*100
```

Out[72]:

81.57894736842105

In [73]:

```

accuracy=confusion_matrix(ytest,predt)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predt, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predt))
print("\n")
print(confusion_matrix(ytest,predt))
print("\n")
print(classification_report(ytest,predt))

```

Accuracy: 81.57894736842105
 Probability of detection of defect(Recall, pd): 0.7647058823529411
 Probability of false alarm(pf): 0.18181818181818182
 Probability of correct detection(Precision): 0.8125

F1-score or FM: 0.787878787878788
 AUC value: 0.8109243697478993

```

[[18  3]
 [ 4 13]]


```

	precision	recall	f1-score	support
0.0	0.82	0.86	0.84	21
1.0	0.81	0.76	0.79	17
accuracy			0.82	38
macro avg	0.82	0.81	0.81	38
weighted avg	0.82	0.82	0.82	38

In [74]:


```
x=np.arange(100)*0.01
y=x

disp = plot_roc_curve(tmodel, xtest, ytest)
plt.plot(x,y, '--')
plt.show()
```



In [75]:

```
disp = plot_precision_recall_curve(tmodel, xtest, ytest)
plt.show()
```



ENSEMBLE PREDICTORS

1- ADABOOST

In [76]:

```
from sklearn.ensemble import AdaBoostClassifier
```

In [77]:

```
adamodel = AdaBoostClassifier(n_estimators=100)
adamodel.fit(xtrain,ytrain)
```

Out[77]:

```
AdaBoostClassifier(algorithm='SAMME.R', base_estimator=None, learning_rate=
1.0,
                    n_estimators=100, random_state=None)
```

In [78]:

```
predada=adamodel.predict(xtest)
adamodel.score(xtest,ytest)*100
```

Out[78]:

```
71.05263157894737
```


In [79]:

```

accuracy=confusion_matrix(ytest,predada)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predada, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predada))
print("\n")
print(confusion_matrix(ytest,predada))
print("\n")
print(classification_report(ytest,predada))

```

```

Accuracy: 71.05263157894737
Probability of detection of defect(Recall, pd): 0.7058823529411765
Probability of false alarm(pf): 0.25
Probability of correct detection(Precision): 0.6666666666666666

```

```

F1-score or FM: 0.6857142857142857
AUC value: 0.7100840336134454

```

```

[[15  6]
 [ 5 12]]


```

	precision	recall	f1-score	support
0.0	0.75	0.71	0.73	21
1.0	0.67	0.71	0.69	17
accuracy			0.71	38
macro avg	0.71	0.71	0.71	38
weighted avg	0.71	0.71	0.71	38

In [80]:


```
x=np.arange(100)*0.01
y=x

disp = plot_roc_curve(adamodel, xtest, ytest)
plt.plot(x,y, '--')
plt.show()
```



In [81]:

```
disp = plot_precision_recall_curve(adamodel, xtest, ytest)
plt.show()
```



2- BAGGING

In [82]:

```
from sklearn.ensemble import BaggingClassifier
```

In [83]:

```
bagmodel = BaggingClassifier(base_estimator=None, n_estimators=10) #default=decision tree,  
bagmodel.fit(xtrain, ytrain)
```

Out[83]:

```
BaggingClassifier(base_estimator=None, bootstrap=True, bootstrap_features=False,  
                  max_features=1.0, max_samples=1.0, n_estimators=10,  
                  n_jobs=None, oob_score=False, random_state=None, verbose=0,  
                  warm_start=False)
```

In [84]:

```
predbag=bagmodel.predict(xtest)  
bagmodel.score(xtest, ytest)*100
```

Out[84]:

```
78.94736842105263
```

In [85]:

```

accuracy=confusion_matrix(ytest,predbag)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predbag, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predbag))
print("\n")
print(confusion_matrix(ytest,predbag))
print("\n")
print(classification_report(ytest,predbag))

```

Accuracy: 78.94736842105263
 Probability of detection of defect(Recall, pd): 0.8235294117647058
 Probability of false alarm(pf): 0.15789473684210525
 Probability of correct detection(Precision): 0.7368421052631579

F1-score or FM: 0.7777777777777778
 AUC value: 0.7927170868347339

```

[[16  5]
 [ 3 14]]


```

	precision	recall	f1-score	support
0.0	0.84	0.76	0.80	21
1.0	0.74	0.82	0.78	17
accuracy			0.79	38
macro avg	0.79	0.79	0.79	38
weighted avg	0.80	0.79	0.79	38

In [86]:

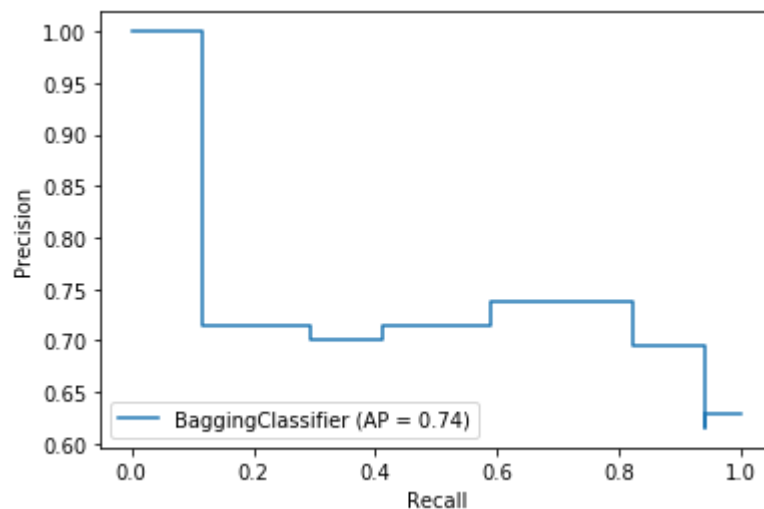
```
x=np.arange(100)*0.01
y=x

disp = plot_roc_curve(bagmodel, xtest, ytest)
plt.plot(x,y, '--')
plt.show()
```



In [87]:

```
disp = plot_precision_recall_curve(bagmodel, xtest, ytest)
plt.show()
```



3- Extra_Tree_Classifier

In [88]:

```
from sklearn.ensemble import ExtraTreesClassifier
```

In [89]:

```
exmodel = ExtraTreesClassifier(n_estimators=100)
exmodel.fit(xtrain, ytrain)
```

Out[89]:

```
ExtraTreesClassifier(bootstrap=False, ccp_alpha=0.0, class_weight=None,
                    criterion='gini', max_depth=None, max_features='auto',
                    max_leaf_nodes=None, max_samples=None,
                    min_impurity_decrease=0.0, min_impurity_split=None,
                    min_samples_leaf=1, min_samples_split=2,
                    min_weight_fraction_leaf=0.0, n_estimators=100,
                    n_jobs=None, oob_score=False, random_state=None, verbose
e=0,
                    warm_start=False)
```

In [90]:

```
predex=exmodel.predict(xtest)
exmodel.score(xtest,ytest)*100
```

Out[90]:

76.31578947368422

In [91]:

```

accuracy=confusion_matrix(ytest,predex)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predex, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predex))
print("\n")
print(confusion_matrix(ytest,predex))
print("\n")
print(classification_report(ytest,predex))

```

Accuracy: 76.31578947368422
 Probability of detection of defect(Recall, pd): 0.6470588235294118
 Probability of false alarm(pf): 0.25
 Probability of correct detection(Precision): 0.7857142857142857

F1-score or FM: 0.7096774193548386
 AUC value: 0.7521008403361344

```

[[18  3]
 [ 6 11]]


```

	precision	recall	f1-score	support
0.0	0.75	0.86	0.80	21
1.0	0.79	0.65	0.71	17
accuracy			0.76	38
macro avg	0.77	0.75	0.75	38
weighted avg	0.77	0.76	0.76	38

In [92]:


```
x=np.arange(100)*0.01
y=x

disp = plot_roc_curve(exmodel, xtest, ytest)
plt.plot(x,y, '--')
plt.show()
```



In [93]:

```
disp = plot_precision_recall_curve(exmodel, xtest, ytest)
plt.show()
```



4- Gradient_Boosting_Classifier

In [94]:

```
from sklearn.ensemble import GradientBoostingClassifier
```


In [95]:

```
gradmodel = GradientBoostingClassifier()  
gradmodel.fit(xtrain,ytrain)
```

Out[95]:

```
GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse', init=None,  
                           learning_rate=0.1, loss='deviance', max_depth=3,  
                           max_features=None, max_leaf_nodes=None,  
                           min_impurity_decrease=0.0, min_impurity_split=None,  
                           min_samples_leaf=1, min_samples_split=2,  
                           min_weight_fraction_leaf=0.0, n_estimators=100,  
                           n_iter_no_change=None, presort='deprecated',  
                           random_state=None, subsample=1.0, tol=0.0001,  
                           validation_fraction=0.1, verbose=0,  
                           warm_start=False)
```

In [96]:

```
predgrad=gradmodel.predict(xtest)  
gradmodel.score(xtest,ytest)*100
```

Out[96]:

71.05263157894737

In [97]:

```

accuracy=confusion_matrix(ytest,predgrad)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predgrad, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predgrad))
print("\n")
print(confusion_matrix(ytest,predgrad))
print("\n")
print(classification_report(ytest,predgrad))

```

Accuracy: 71.05263157894737
 Probability of detection of defect(Recall, pd): 0.7058823529411765
 Probability of false alarm(pf): 0.25
 Probability of correct detection(Precision): 0.6666666666666666

F1-score or FM: 0.6857142857142857
 AUC value: 0.7100840336134454

```


[[15  6]
 [ 5 12]]

```

	precision	recall	f1-score	support
0.0	0.75	0.71	0.73	21
1.0	0.67	0.71	0.69	17
accuracy			0.71	38
macro avg	0.71	0.71	0.71	38
weighted avg	0.71	0.71	0.71	38


In [98]:

```
x=np.arange(100)*0.01  
y=x  
  
disp = plot_roc_curve(gradmodel, xtest, ytest)  
plt.plot(x,y, '--')  
plt.show()
```



In [99]:

```
disp = plot_precision_recall_curve(gradmodel, xtest, ytest)  
plt.show()
```



5- Random_Forest_Classifier

In [100]:

```
from sklearn.ensemble import RandomForestClassifier
```

In [101]:

```
randmodel = RandomForestClassifier()  
randmodel.fit(xtrain,ytrain)
```

Out[101]:

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,  
                        criterion='gini', max_depth=None, max_features='auto',  
                        max_leaf_nodes=None, max_samples=None,  
                        min_impurity_decrease=0.0, min_impurity_split=None,  
                        min_samples_leaf=1, min_samples_split=2,  
                        min_weight_fraction_leaf=0.0, n_estimators=100,  
                        n_jobs=None, oob_score=False, random_state=None,  
                        verbose=0, warm_start=False)
```

In [102]:

```
predrand=randmodel.predict(xtest)  
randmodel.score(xtest,ytest)*100
```

Out[102]:

78.94736842105263

In [103]:

```

accuracy=confusion_matrix(ytest,predrand)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predrand, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predrand))
print("\n")
print(confusion_matrix(ytest,predrand))
print("\n")
print(classification_report(ytest,predrand))

```

Accuracy: 78.94736842105263
 Probability of detection of defect(Recall, pd): 0.8235294117647058
 Probability of false alarm(pf): 0.15789473684210525
 Probability of correct detection(Precision): 0.7368421052631579

F1-score or FM: 0.7777777777777778
 AUC value: 0.7927170868347339

```

[[16  5]
 [ 3 14]]


```

	precision	recall	f1-score	support
0.0	0.84	0.76	0.80	21
1.0	0.74	0.82	0.78	17
accuracy			0.79	38
macro avg	0.79	0.79	0.79	38
weighted avg	0.80	0.79	0.79	38

In [104]:


```
x=np.arange(100)*0.01
y=x

disp = plot_roc_curve(randmodel, xtest, ytest)
plt.plot(x,y, '--')
plt.show()
```



In [105]:

```
disp = plot_precision_recall_curve(randmodel, xtest, ytest)
plt.show()
```



6- Stacking_Classifier

In [106]:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import LinearSVC
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.ensemble import StackingClassifier
```

In [107]:

```

estimators = [('rf', RandomForestClassifier(n_estimators=10, random_state=42)),
              ('svr', make_pipeline(StandardScaler(), LinearSVC(random_state=42)))]

stmodel = StackingClassifier(estimators=estimators, final_estimator=LogisticRegression())
stmodel.fit(xtrain,ytrain)

```

```

C:\ProgramData\Anaconda3\envs\myenv\lib\site-packages\sklearn\svm\_base.py:9
47: ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
    "the number of iterations.", ConvergenceWarning)
C:\ProgramData\Anaconda3\envs\myenv\lib\site-packages\sklearn\svm\_base.py:9
47: ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
    "the number of iterations.", ConvergenceWarning)
C:\ProgramData\Anaconda3\envs\myenv\lib\site-packages\sklearn\svm\_base.py:9
47: ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
    "the number of iterations.", ConvergenceWarning)
C:\ProgramData\Anaconda3\envs\myenv\lib\site-packages\sklearn\svm\_base.py:9
47: ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
    "the number of iterations.", ConvergenceWarning)
C:\ProgramData\Anaconda3\envs\myenv\lib\site-packages\sklearn\svm\_base.py:9
47: ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
    "the number of iterations.", ConvergenceWarning)
C:\ProgramData\Anaconda3\envs\myenv\lib\site-packages\sklearn\svm\_base.py:9
47: ConvergenceWarning: Liblinear failed to converge, increase the number of
iterations.
    "the number of iterations.", ConvergenceWarning)

```

Out[107]:

```

StackingClassifier(cv=None,
                  estimators=[('rf',
                              RandomForestClassifier(bootstrap=True,
                                                         ccp_alpha=0.0,
                                                         class_weight=None,
                                                         criterion='gini',
                                                         max_depth=None,
                                                         max_features='aut
o',
                                                         max_leaf_nodes=None,
                                                         max_samples=None,
                                                         min_impurity_decrea
se=0.0,
                                                         min_impurity_split=
None,
                                                         min_samples_leaf=1,
                                                         min_samples_split=
2,
                                                         min_weight_fraction
_leaf=0.0,
                                                         n_estimators=10,
                                                         n_jobs=None,...

```

```

                                tol=0.0001,
                                verbose=0))],
                                verbose=False))],
final_estimator=LogisticRegression(C=1.0, class_weight=
None,
                                dual=False,
                                fit_intercept=True,
                                intercept_scaling=1,
                                l1_ratio=None,
                                max_iter=100,
                                multi_class='auto',
                                n_jobs=None, penalty
='l2',
                                random_state=None,
                                solver='lbfgs',
                                tol=0.0001, verbose=
0,
                                warm_start=False),
n_jobs=None, passthrough=False, stack_method='auto',
verbose=0)

```

In [108]:

```

predst=stmodel.predict(xtest)
stmodel.score(xtest,ytest)*100

```

Out[108]:

81.57894736842105

In [109]:

```

accuracy=confusion_matrix(ytest,predst)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predst, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predst))
print("\n")
print(confusion_matrix(ytest,predst))
print("\n")
print(classification_report(ytest,predst))

```

```

Accuracy: 81.57894736842105
Probability of detection of defect(Recall, pd): 0.8823529411764706
Probability of false alarm(pf): 0.11111111111111111
Probability of correct detection(Precision): 0.75

```

```

F1-score or FM: 0.8108108108108107
AUC value: 0.8221288515406161

```

```

[[16  5]
 [ 2 15]]


```

	precision	recall	f1-score	support
0.0	0.89	0.76	0.82	21
1.0	0.75	0.88	0.81	17
accuracy			0.82	38
macro avg	0.82	0.82	0.82	38
weighted avg	0.83	0.82	0.82	38

In [110]:

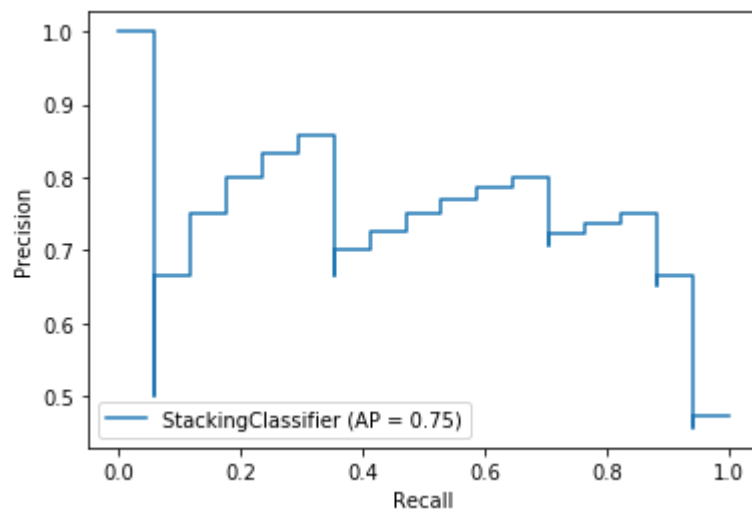
```
x=np.arange(100)*0.01
y=x

disp = plot_roc_curve(stmodel, xtest, ytest)
plt.plot(x,y, '--')
plt.show()
```



In [111]:

```
disp = plot_precision_recall_curve(stmodel, xtest, ytest)
plt.show()
```



7- Voting_Classifier

In [112]:

```
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
```

In [113]:

```

clf1 = LogisticRegression()
clf2 = RandomForestClassifier(#n_estimators=50, random_state=1)
clf3 = GaussianNB()
votmodel = VotingClassifier(estimators=[('lr', clf1), ('rf', clf2), ('gnb', clf3)], voting=
votmodel.fit(xtrain,ytrain)

```

Out[113]:

```

VotingClassifier(estimators=[('lr',
                             LogisticRegression(C=1.0, class_weight=None,
                                                    dual=False, fit_intercept=T
rue,
                                                    intercept_scaling=1,
                                                    l1_ratio=None, max_iter=10
0,
                                                    multi_class='auto',
                                                    n_jobs=None, penalty='l2',
                                                    random_state=None,
                                                    solver='lbfgs', tol=0.0001,
                                                    verbose=0, warm_start=False
e)),
                  ('rf',
                   RandomForestClassifier(bootstrap=True,
                                          ccp_alpha=0.0,
                                          class_weight=None,
                                          cr...
                                          max_leaf_nodes=None,
                                          max_samples=None,
                                          min_impurity_decrease=
0.0,
                                          min_impurity_split=None
e,
                                          min_samples_leaf=1,
                                          min_samples_split=2,
                                          min_weight_fraction_lea
f=0.0,
                                          n_estimators=100,
                                          n_jobs=None,
                                          oob_score=False,
                                          random_state=None,
                                          verbose=0,
                                          warm_start=False)),
                  ('gnb',
                   GaussianNB(priors=None, var_smoothing=1e-0
9)]],
              flatten_transform=True, n_jobs=None, voting='hard',
              weights=None)

```

In [114]:

```

predvot=votmodel.predict(xtest)
votmodel.score(xtest,ytest)*100

```

Out[114]:

73.68421052631578

In [115]:

```

accuracy=confusion_matrix(ytest,predvot)
TP=accuracy[0][0]
FP=accuracy[1][0]
TN=accuracy[1][1]
FN=accuracy[0][1]

print("Accuracy: ",(TP+TN)/(TP+FP+TN+FN)*100)
print("Probability of detection of defect(Recall, pd): ",TN/(TN+FP))
print("Probability of false alarm(pf): ",FP/(TP+FP))
print("Probability of correct detection(Precision): ", TN/(TN+FN))
print("\n")
print("F1-score or FM: ", f1_score(ytest, predvot, average='binary'))
print("AUC value: ",roc_auc_score(ytest, predvot))
print("\n")
print(confusion_matrix(ytest,predvot))
print("\n")
print(classification_report(ytest,predvot))

```

Accuracy: 73.68421052631578
 Probability of detection of defect(Recall, pd): 1.0
 Probability of false alarm(pf): 0.0
 Probability of correct detection(Precision): 0.6296296296296297

F1-score or FM: 0.7727272727272727
 AUC value: 0.7619047619047619

```

[[11 10]
 [ 0 17]]

```

	precision	recall	f1-score	support
0.0	1.00	0.52	0.69	21
1.0	0.63	1.00	0.77	17
accuracy			0.74	38
macro avg	0.81	0.76	0.73	38
weighted avg	0.83	0.74	0.73	38

In []:

In []: